

```
library(dplyr)

rladies_global %>%
  filter(city == 'Bari')
```



Manipolare i dati con Tidyverse





Hello!

Welcome to R-Ladies
Bari



1. Introduction

R language, RStudio



Three things you'll need to install

1.

Install R -- this is the open-source programming language we'll use (download via CRAN -- Comprehensive R Archive Network)

2.

Install RStudio -- this is the most popular IDE for R and will make your life a lot easier (download from rstudio.com/download)

3.

Install the tidyverse -- this is the group of packages we'll use within R to work with data.

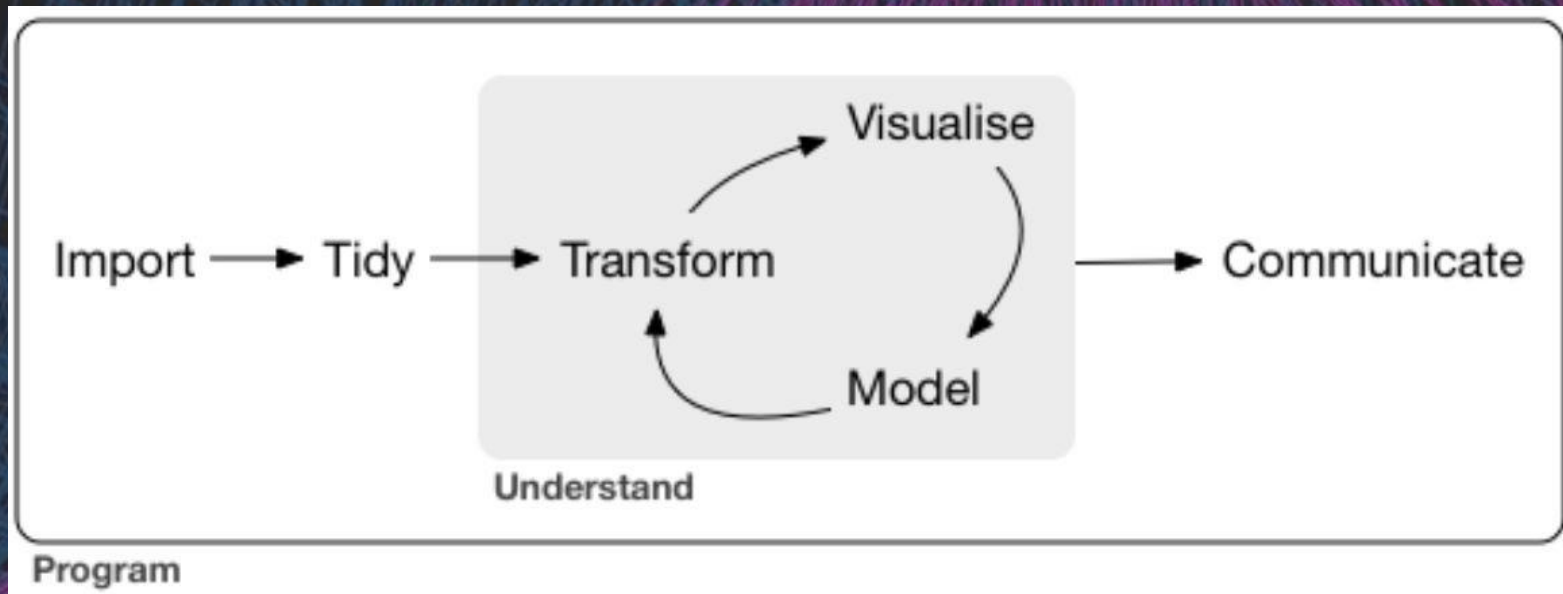
Install with one line of code in R: `install.packages("tidyverse")`



2. Manipolare i dati con Tidyverse



The data science process (tidied)

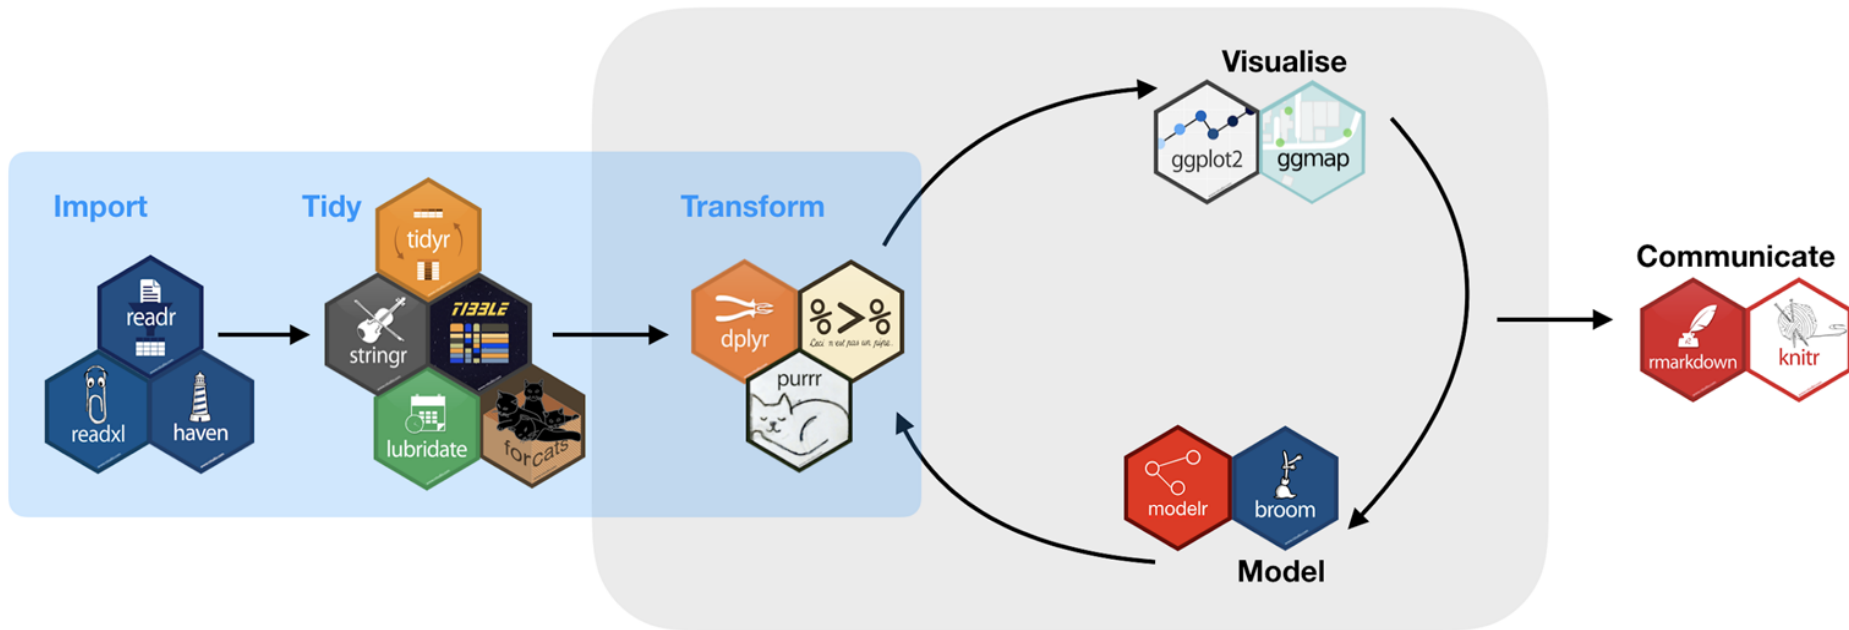




What is the tidyverse?

- **Collection of R packages** based on tidy data principles
- Designed to **work together**
- An **easier** way to code!
- AKA “Hadleyverse” (most packages written by Hadley Wickham)

The data science process with tidyverse



What is tidy data?

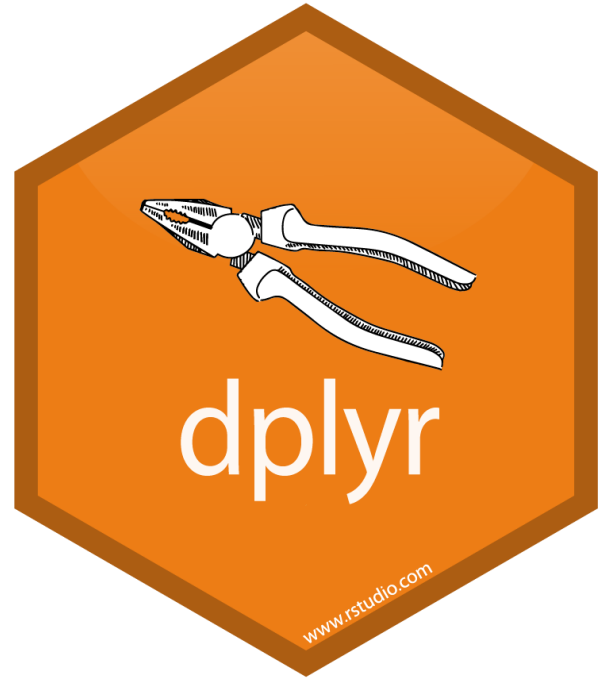
- + Each **variable** is a column
- + Each observation is a row
- + Each type of observational unit is a table

id	artist	track	time
1	2 Pac	Baby Don't Cry	4:22
2	2Ge+her	The Hardest Part Of ...	3:15
3	3 Doors Down	Kryptonite	3:53
4	3 Doors Down	Loser	4:24
5	504 Boyz	Wobble Wobble	3:35
6	98~0	Give Me Just One Nig...	3:24
7	A*Teens	Dancing Queen	3:44
8	Aaliyah	I Don't Wanna	4:15
9	Aaliyah	Try Again	4:03
10	Adams, Yolanda	Open My Heart	5:30
11	Adkins, Trace	More	3:05
12	Aguilera, Christina	Come On Over Baby	3:38
13	Aguilera, Christina	I Turn To You	4:00
14	Aguilera, Christina	What A Girl Wants	3:18
15	Alice DeeJay	Better Off Alone	6:50



3. dplyr

Let's start with the first set of slides



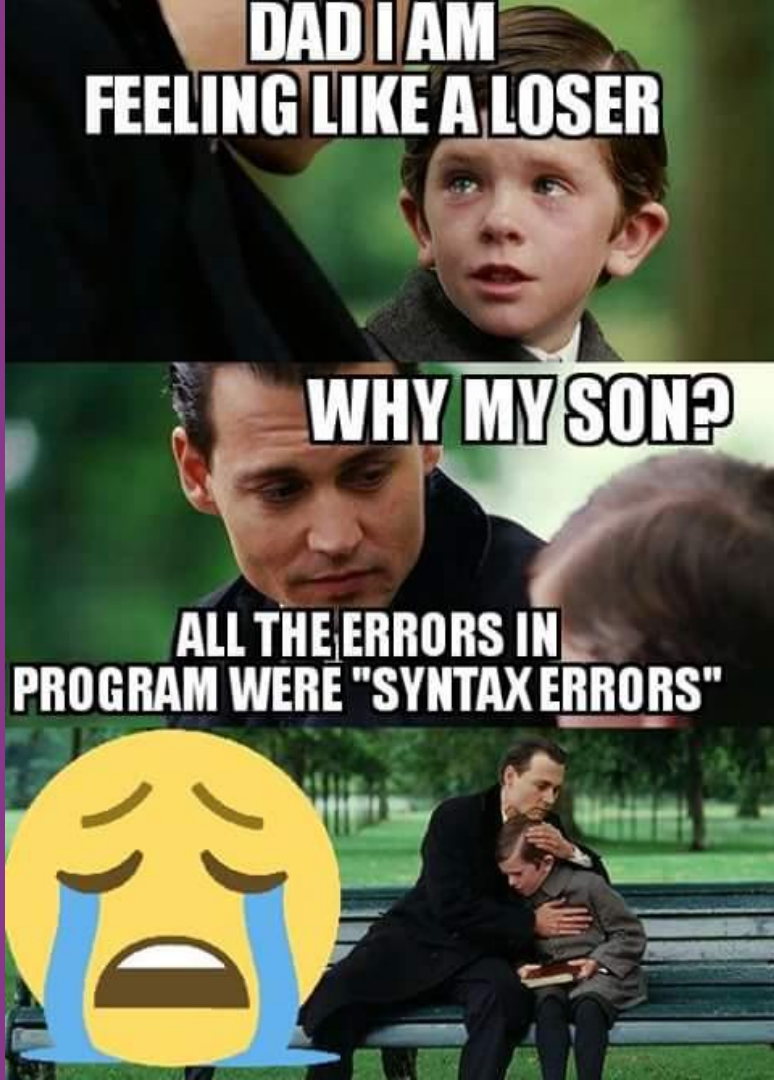


dplyr



grammar of data manipulation

- + `mutate()` to create new variables from existing ones
- + `select()` picks variables based on their names
- + `filter()` allows pointed selection based on given criteria
- + `summarise()` reduces multiple values down to a single summary
- + `arrange()` changes the ordering of rows
- + `group_by()` performs any of the above on a group-by-group basis



dplyr syntax



- + All **calls** to dplyr verbs follow the same **format**:
 1. The first argument is a **dataframe**
 2. The subsequent arguments describe **what to do** to that dataframe, using unquoted variable names.
- + Each call returns a new dataframe (rather than overwriting the 'old' one)
- + Example:

```
filter(dataset_name,  
       name == "Francesca")
```

What is magrittr?

Simplifying R code with **pipes** ($\%>\%$)

- + Easy way to pass data through functions without nesting
- + First argument of each function is “piped” in to reduce redundancy



dplyr + magrittr

example

before

```
summarise(  
  group_by((filter  
    (babynames, name == "Caitlin")  
  ),  
    year  
  ),  
  total = sum(n))
```

after

```
babynames %>%  
  filter(name == "Caitlin") %>%  
  group_by(year) %>%  
  summarise(total = sum(n))
```


Quick aside: iris dataset



- Included in R (**iris** to view)
- 150 observations of 5 variables:
iris type, sepal length, sepal width, petal
length and petal width



select()

- + Picks variables based on their names
- + First argument is dataframe; subsequent arguments represent columns to select



```
iris %>% select(Species, Petal.Length, Petal.Width)
```

```
> iris
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1         5.1         3.5          1.4          0.2   setosa
2         4.9         3.0          1.4          0.2   setosa
3         4.7         3.2          1.3          0.2   setosa
4         4.6         3.1          1.5          0.2   setosa
5         5.0         3.6          1.4          0.2   setosa
```

```
Species Petal.Length Petal.Width
1   setosa          1.4          0.2
2   setosa          1.4          0.2
3   setosa          1.3          0.2
4   setosa          1.5          0.2
5   setosa          1.4          0.2
```


select() + helper functions

Helper functions you can use within **select()**:

- + **starts_with("a")** matches names that begin with “a”
- + **ends_with("z")** matches names that begin with “z”
- + **contains("lady")** matches names that contain “lady”
- + **matches(<regex>)** allows you to do regex matching on names

```
> iris %>% select(starts_with("p"))  
  Petal.Length Petal.Width  
1          1.4          0.2  
2          1.4          0.2  
3          1.3          0.2  
4          1.5          0.2  
5          1.4          0.2
```

filter()

- + Allows pointed selection based on given criteria
- + First argument is the dataframe, subsequent arguments are expressions used to filter the dataframe



```
iris %>% filter(Species == "setosa")
```

```
> iris
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1         3.5          1.4          0.2  setosa
2          4.9         3.0          1.4          0.2  setosa
3          4.7         3.2          1.3          0.2  setosa
4          4.6         3.1          1.5          0.2  setosa
5          5.0         3.6          1.4          0.2  setosa
```

nrow = 150

```
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1         3.5          1.4          0.2  setosa
2          4.9         3.0          1.4          0.2  setosa
3          4.7         3.2          1.3          0.2  setosa
4          4.6         3.1          1.5          0.2  setosa
5          5.0         3.6          1.4          0.2  setosa
```

nrow = 50

filter() booleans

+



Multiple arguments to `filter()` are combined with “and”: every expression must be true in order for a row to be included in the output. For other types of combinations, you’ll need to use Boolean operators yourself: `&` is “and”, `|` is “or”, and `!` is “not”. Figure 5.1 shows the complete set of Boolean operations.

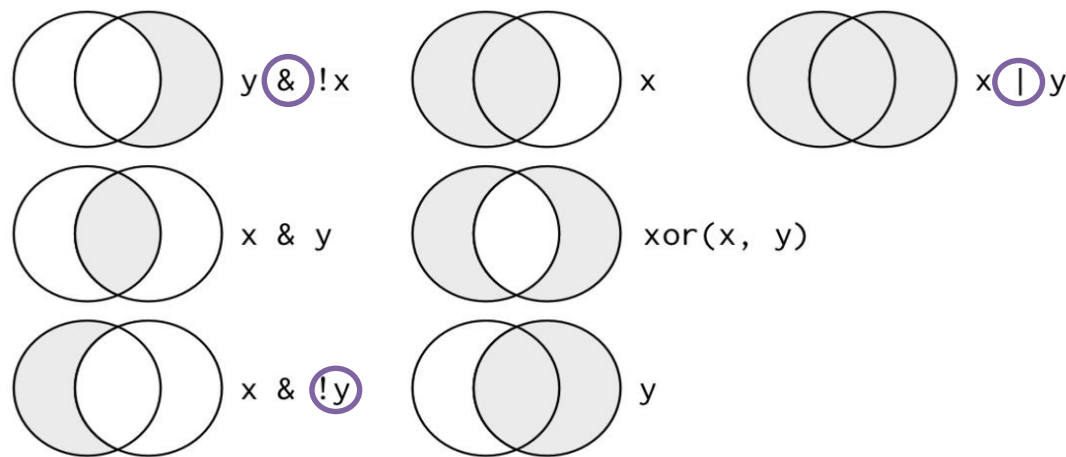


Figure 5.1: Complete set of boolean operations. x is the left-hand circle, y is the right-hand circle, and the shaded region show which parts each operator selects.

arrange()

- + Changes the ordering of rows
- + First argument is the dataframe, subsequent arguments are columns and/or expressions used to re-arrange the dataframe
- + Note: default is ascending order, and NA's are always at the end



```
iris %>% arrange(Sepal.Length, Sepal.Width)
```

```
> iris
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1         5.1         3.5          1.4          0.2   setosa
2         4.9         3.0          1.4          0.2   setosa
3         4.7         3.2          1.3          0.2   setosa
4         4.6         3.1          1.5          0.2   setosa
5         5.0         3.6          1.4          0.2   setosa
```

```
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1         4.3         3.0          1.1          0.1   setosa
2         4.4         2.9          1.4          0.2   setosa
3         4.4         3.0          1.3          0.2   setosa
4         4.4         3.2          1.3          0.2   setosa
5         4.5         2.3          1.3          0.3   setosa
```

Quick aside: Missing values

- + **NA** represents a *missing (unknown)* value
- + Comparisons involve unknown values typically result in unknown values
- + To see whether a value is missing, use **is.na()**
- + **filter()** only includes rows where the condition is **true** (not false or **NA**)

```
# Let x be Mary's age. We don't know how old she is.  
x <- NA  
  
# Let y be John's age. We don't know how old he is.  
y <- NA  
  
# Are John and Mary the same age?  
x == y  
#> [1] NA  
# We don't know!
```



mutate()



- + Creates new variables from existing ones
- + Note: columns created with **mutate()** are always added to end of dataset

```
iris %>% mutate(petal_area = Petal.Length * Petal.Width)
```

```
> iris
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1          3.5          1.4          0.2  setosa
2          4.9          3.0          1.4          0.2  setosa
3          4.7          3.2          1.3          0.2  setosa
4          4.6          3.1          1.5          0.2  setosa
5          5.0          3.6          1.4          0.2  setosa
```

```
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species petal_area
1          5.1          3.5          1.4          0.2  setosa      0.28
2          4.9          3.0          1.4          0.2  setosa      0.28
3          4.7          3.2          1.3          0.2  setosa      0.26
4          4.6          3.1          1.5          0.2  setosa      0.30
5          5.0          3.6          1.4          0.2  setosa      0.28
```

mutate()

Useful functions

- + Arithmetic operators (+, -, *, /, ^)
- + Log functions (like `log10()`)
- + Offsets like `lead()` and `lag()`
- + Logical comparisons (<, <=, >, >=, !=)
- + `ifelse()` statements (*if this, then this, else this*)
- + Cumulative and rolling aggregates
- + Ranking (like `ntile()`)

group_by() and summarise()

- + group_by applies dplyr verbs by group
- + summarise reduces multiple values down to a single summary

```
iris %>%  
  group_by(Species) %>%  
  summarise(avg_petal_width = mean(Petal.Width))
```

```
> iris
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa

```
# A tibble: 3 × 2  
  Species avg_petal_width  
  <fctr>      <dbl>  
1  setosa      0.246  
2 versicolor  1.326  
3  virginica   2.026
```


summarise()

Useful functions

- + Counts (`n()`, `n_distinct()`)
- + Measures of location (`mean()`, `median()`)
- + Measures of spread (`sd()`, `IQR()`)
- + Measures of position (`first()`, `last()`)

Tips & Tricks

- + If you don't have the result of a dplyr chain to a dataframe, it will print
- + If you want to print and save, wrap assignment in parenthesis
Example: `(iris_names <- iris %>% filter(Species == "setosa"))`
- + `rename()` is a cool function to clean up messy column names
- + After grouping with `group_by()`, you can `ungroup()` to remove groupings
- + There is a [cheat sheet](#) for data wrangling!

Cheat sheet for data wrangling

Data Transformation with dplyr : CHEAT SHEET



dplyr functions work with pipes and expect tidy data. In tidy data:



Each **variable** is in its own **column**



Each **observation**, or **case**, is in its own **row**



$x \%>\% f(y)$ becomes $f(x, y)$

pipes

Summarise Cases

These apply **summary functions** to columns to create a new table of summary statistics. Summary functions take vectors as input and return one value (see back).

summary function



summarise(data, ...) Compute table of summaries. `summarise(mtcars, avg = mean(mpg))`



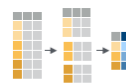
count(x, ..., wt = NULL, sort = FALSE) Count number of rows in each group defined by the variables in ... Also **tally**(). `count(iris, Species)`

VARIATIONS

summarise_all() - Apply funs to every column.
summarise_at() - Apply funs to specific columns.
summarise_if() - Apply funs to all cols of one type.

Group Cases

Use **group_by()** to create a "grouped" copy of a table. dplyr functions will manipulate each "group" separately and then combine the results.



`mtcars %>%
group_by(cyl) %>%
summarise(avg = mean(mpg))`

group_by(data, ..., add = FALSE)
Returns copy of table grouped by ...
`g_iris <- group_by(iris, Species)`

ungroup(x, ...) Returns ungrouped copy of table.
`ungroup(g_iris)`

Manipulate Cases

EXTRACT CASES

Row functions return a subset of rows as a new table.



filter(data, ...) Extract rows that meet logical criteria. `filter(iris, Sepal.Length > 7)`



distinct(data, ..., keep_all = FALSE) Remove rows with duplicate values. `distinct(iris, Species)`



sample_frac(tbl, size = 1, replace = FALSE, weight = NULL, env = parent.frame()) Randomly select fraction of rows.
`sample_frac(iris, 0.5, replace = TRUE)`



sample_n(tbl, size, replace = FALSE, weight = NULL, env = parent.frame()) Randomly select size rows. `sample_n(iris, 10, replace = TRUE)`



slice(data, ...) Select rows by position. `slice(iris, 10:15)`

top_n(x, n, wt) Select and order top n entries (by group if grouped data). `top_n(iris, 5, Sepal.Width)`

Logical and boolean operators to use with filter()

`<` `<=` `is.na()` `%in%` `|` `xor()`
`>` `>=` `!is.na()` `!` `&`

See ?base::logic and ?Comparison for help.

ARRANGE CASES



arrange(data, ...) Order rows by values of a column or columns (low to high), use with **desc()** to order from high to low.

`arrange(mtcars, mpg)`
`arrange(mtcars, desc(mpg))`

ADD CASES



add_row(data, ..., before = NULL, after = NULL) Add one or more rows to a table.
`add_row(faithful, eruptions = 1, waiting = 1)`

Manipulate Variables

EXTRACT VARIABLES

Column functions return a set of columns as a new vector or table.



pull(data, var = -1) Extract column values as a vector. Choose by name or index. `pull(iris, Sepal.Length)`



select(data, ...) Extract columns as a table. Also **select_if()**. `select(iris, Sepal.Length, Species)`

Use these helpers with **select()**, e.g. `select(iris, starts_with("Sepal"))`

contains(match) **num_range**(prefix, range) ; e.g. `mpg:cyl`
ends_with(match) **one_of**(...) ; e.g. `Species`
matches(match) **starts_with**(match)

MAKE NEW VARIABLES

These apply **vectorized functions** to columns. Vectorized functions take vectors as input and return vectors of the same length as output (see back).

vectorized function



mutate(data, ...) Compute new column(s). `mutate(mtcars, gpm = 1/mpg)`



transmute(data, ...) Compute new column(s), drop others. `transmute(mtcars, gpm = 1/mpg)`



mutate_all(tbl, funs, ...) Apply funs to every column. Use with **funs()**. Also **mutate_if()**. `mutate_all(faithful, funs(log(), log2(), log10()))`
`mutate_if(iris, is.numeric, funs(log(), log10()))`



mutate_at(tbl, cols, funs, ...) Apply funs to specific columns. Use with **funs()**, **vars()** and the helper functions for **select()**. `mutate_at(iris, vars(-Species), funs(log(), log10()))`



add_column(data, ..., before = NULL, after = NULL) Add new column(s). Also **add_count()**, **add_tally()**. `add_column(mtcars, new = 1:32)`



rename(data, ...) Rename columns. `rename(iris, Length = Sepal.Length)`



4. Wrap-up

Announcements, upcoming events, etc.

R-Ladies Bari

Upcoming Events





R-Ladies Bari

E-Learning

Esercitazioni gratuite

<https://www.datacamp.com/courses/introduction-to-the-tidyverse>

<https://www.kaggle.com/rtatman/manipulating-data-with-the-tidyverse>